

AppSec Platform Blueprint

System Architecture

Our AppSec platform will adopt a microservices architecture to maximize scalability and resilience. Each function (e.g. SAST engine, SCA engine, results database, UI/API gateway) runs as an independent service in containers orchestrated by Kubernetes ¹. For example, Atlassian notes that “a microservices architecture splits an application into a series of independently deployable services” ¹, each of which can be scaled or updated without affecting others. We will host core control planes (e.g. orchestration, auth, analytics) in the cloud and allow the data plane (scanners, agents) to run on-premises or in customers’ private clouds. Code commits or pull requests trigger scanning jobs: a job queue (e.g. Kafka or Kubernetes Jobs) dispatches each project to the appropriate SAST or SCA scanner service. Results are normalized (e.g. into SARIF) and stored in a central database (SQL or document store) for querying and reporting. Integrations (webhooks, API endpoints) connect with CI/CD pipelines and issue trackers for automatic triage. *Figure: A sample microservices architecture (Atlassian example) where individual services (account, inventory, cart, etc.) communicate via APIs; this decoupling improves scalability and maintainability* ¹.




Figure: Example microservices architecture (source: Atlassian ¹). Each service handles a specific function, communicating via APIs and an API gateway.

Our platform core will include: an API gateway (routing and authentication), user interface/dashboard, pipeline orchestration service, SAST orchestration service, SCA/SBOM service, vulnerability database/updater, AI assistant service, and logging/metrics services. Communication between components will use secured gRPC or REST interfaces. Scan jobs will run in ephemeral worker containers: for SAST, workers load source code, run analyzers (like Semgrep, CodeQL, or custom LLM checks), and emit findings. For SCA, workers generate SBOMs (e.g. via Syft) and query vulnerability feeds. A policy engine (e.g. Open Policy

Agent) enforces access controls and enforces compliance checks at each stage. High availability is ensured by deploying services across multiple nodes or zones. We will standardize on container images and manifest definitions (Helm/Terraform) so the platform can be deployed in any Kubernetes-compatible environment (cloud or on-prem). By splitting components and adopting containers, the design scales with load and enables independent upgrades of scanners, databases, and UIs.

Tech Stack Recommendations

We recommend a **containerized, polyglot microservices stack** using industry-standard tools. For compute, use **Kubernetes** (K8s) on cloud (EKS/GKE/AKS) and on-prem (Rancher/OpenShift) to orchestrate Dockerized services. For programming languages, choose a mix to leverage strengths: for high-performance back-end services, languages like **Go** or **Java/Kotlin** are ideal (Go for lightweight concurrency and small images ² ³, Java/Kotlin for mature enterprise frameworks like Spring Boot). **Python** or **Node.js/TypeScript** can be used for glue logic, AI integration, or web services where rapid development is needed. A **React** (or Angular) front-end communicates via REST or GraphQL.

Key tech suggestions:

- **Scanners:** Use open-source analyzers (Semgrep for many languages, CodeQL for deep Java/C# analysis, SpotBugs/PMD for Java, ESLint/TSLint for JS, Bandit for Python, etc.) invoked via the orchestration service.

- **Message broker:** Use Kafka or RabbitMQ for job dispatching and event streaming.
- **Database:** A relational DB (PostgreSQL) for structured data (user accounts, policies, results meta) plus a document/search store (e.g. Elasticsearch) for fast querying of findings. For SBOM or dependency graphs, consider a graph database (Neo4j or Dgraph) to model component relationships.
- **CI/CD and DevOps:** The platform itself is delivered via Terraform/Helm charts. Use CI pipelines for platform code (GitHub Actions, Jenkins, GitLab CI). For scanning integration, provide native connectors/plugins for GitHub, GitLab, Jenkins, Azure Pipelines, etc. Dockerize all components for portability.
- **Logging/Monitoring:** Use Prometheus/Grafana for metrics, and ELK/EFK (Elasticsearch, Fluentd, Kibana) or a managed service (Splunk, CloudWatch) for logs and audit trails.
- **Auth & Security:** Employ OAuth2/OpenID (e.g. Keycloak or Auth0) for single sign-on and RBAC. Encrypt all data in transit (TLS 1.3) and at rest (AES-256).
- **AI Infrastructure:** For AI/ML, leverage cloud services or local servers with GPUs. Use containerized model serving (e.g. Hugging Face's `transformers` in a REST API) for any self-hosted models. Models like GPT-4 can be accessed via API, and CodeLlama or Llama 3 can be hosted via open-source runtimes. For vector similarity (e.g. code embeddings), a vector DB (Milvus, Pinecone) could be added. Use Python notebooks or MLflow for any custom model training pipelines.

For long-term scalability, avoid monolith frameworks. Use asynchronous processing (job queues) so a failing scanner doesn't block others. Apply 12-Factor App principles (externalize config, stateless services, logs as event streams ¹). Over time, this stack can evolve: e.g., new languages can be supported by adding new scanner containers, and new cloud providers by writing new IaC scripts.

AI Model Strategy

The platform will **leverage pre-trained LLMs** (e.g. GPT-4, Meta Code Llama) for tasks like vulnerability detection, fix suggestion, and developer Q&A, with optional fine-tuning on security data. Initially, we can use GPT-4 (via API) for context-aware analysis and use open models (Llama-3/Code Llama) for code-specific

tasks if on-prem or cost is an issue. We will design prompts to clearly delineate tasks (e.g. “Analyze this code snippet for buffer overflow vulnerabilities and explain” or “Suggest a fix for the identified XSS issue in the code below”). To improve accuracy and reduce hallucinations, we may fine-tune smaller models on curated vulnerability datasets. For example, an innovation team fine-tuned Llama-3 on a labeled C/C++ vulnerability dataset and significantly outperformed plain GPT-4 on that task ⁴. Our fine-tuning pipeline can follow a standard process: collect labeled code samples (e.g. from VDISC, Juliet, SARD), preprocess into function-level examples, apply LoRA/QLoRA training, and validate against a held-out test set. *Figure: A typical LLM fine-tuning workflow. A pre-trained model and a labeled code dataset (with vulnerability labels) are used in a supervised fine-tuning step. The fine-tuned model is then deployed for inference (code scanning) and evaluated for accuracy* ⁴.




Figure: Example pipeline for fine-tuning an LLM on code vulnerability data (adapted from inovex GmbH) ⁴.

In practice, we will:

- **Prompt engineering:** Develop structured prompts that include task instructions and examples. Use chain-of-thought prompts if needed (e.g. “Here is code... break down the data flow and identify any unsafe uses”).
- **Retrieval augmentation:** Maintain a knowledge base (e.g. OWASP/CWE descriptions, past findings) to feed relevant context to the model via retrieval.
- **Fine-tuning:** Set up pipelines (e.g. Hugging Face Trainer) to fine-tune open models for tasks like “vulnerability classification (binary)” or “generate secure code snippet given vulnerable code.” Ensure data is diverse (multiple languages, CWEs).
- **Human-in-loop:** Collect user feedback (e.g. “Was this suggestion correct?”) to iteratively improve model prompts or training data.
- **Bias and risk management:** Test models for common failure modes. The literature shows GPT-4 can detect many CWE types (94% accuracy over 32 types) ⁵, but performance can vary by vulnerability class, so LLM suggestions will be treated as high-confidence only when validated (e.g. accompanied by a confidence score or proof).

We will *not* rely on AI alone: as the GPT-4 study notes, “LLMs should augment, not replace, existing security practices” ⁶. Thus, all LLM-generated findings will be cross-checked by the standard SAST/SCA pipeline or

flagged for human review. Over time, an AI “developer assistant” could be added as an IDE plugin, offering inline fixes (similar to Checkmarx Assist), but MVP will focus on backend analysis and dashboard guidance.

SAST Module Design

The SAST subsystem orchestrates multiple static analyzers and unifies their output. The orchestration service will: take a code repository and selected languages, spawn language-specific scanner jobs (e.g. one job runs Semgrep for Python/JS, another runs CodeQL for Java, another runs a custom LLM-based checker for business logic, etc.), then collect all findings. Each scanner should output results in a standardized format. We will normalize these into SARIF (Static Analysis Results Interchange Format) ⁷ because SARIF is an industry-standard JSON schema “designed to facilitate interoperability” between static analysis tools ⁷. Normalizing to SARIF lets us aggregate and de-duplicate findings from different tools easily.

We will support major languages from day one (Java, Python, JavaScript/TypeScript, C#, Go, C/C++, etc.), using appropriate open-source analyzers for each. Custom rules can be added to Semgrep and CodeQL libraries for specific frameworks or in-house patterns. We will implement scan modes: a **fast scan** (checks quick-to-run rules for instant feedback in PRs) and a **full scan** (deep dataflow and all rules for nightly builds). This mirrors best practices: fast scans catch routine issues early in CI/CD, and deeper scans run pre-release ⁸.

To reduce noise, we will incorporate contextual filtering. For example, findings in test code or third-party libraries (if protected by SCA) can be auto-muted. Over time, we may use ML to predict false positives (e.g. pattern-matching models). Regardless, all findings will include metadata (file, line, CWE classification, CVSS score estimate, snippet) and links to references. We will also support incremental scans (only changed code) to accelerate feedback. Overall, by orchestrating SAST engines and standardizing results (using SARIF ⁷, CVSS scores, and CWE tags), we provide a unified view of code-based vulnerabilities across languages and projects.

SCA Module Design

The SCA subsystem continuously inventories open-source components and identifies known vulnerabilities. We will generate a Software Bill of Materials (SBOM) for every project, ideally during the build process ⁹. Using tools like Syft or CycloneDX CLI, we capture all direct and transitive dependencies in a standard format (e.g. CycloneDX JSON or SPDX) ⁹. The SBOM will include component identifiers (purls), versions, checksums, and dependency relationships. We will store and version these SBOMs in a central registry (or use a dedicated SBOM manager like Dependency-Track). This ensures traceability and reproducibility ¹⁰.

For each component in the SBOM, we will query vulnerability databases (e.g. NVD, OSS Index, GitHub Advisory DB). Tools like Grype or OWASP Dependency-Check can map CVEs to SBOM entries ¹¹. We will enrich components with CVE data (severity, descriptions, fixed versions). This allows us to present to users the full inventory of vulnerable libraries, their severity, and remediation advice. We will also track licenses and flag any license compliance issues.

Critically, we will implement **reachability analysis** in SCA: determining if a vulnerable package’s code is actually used by the application. As Labrador Labs explains, modern applications include many unused components, and a vulnerability only matters if the app can execute the vulnerable code ¹². We will

statically analyze dependency call graphs (e.g. using CodeQL or Sonatype's DepGraph) to see if application code ever invokes the vulnerable functions. This way, we can de-prioritize "theoretical" issues. OWASP guidelines encourage this deeper contextual understanding ("look inside the system, understand execution paths" ¹³).

For remediation, we will use the SBOM to triage: mapping CVEs to components, using any available VEX (Vulnerability Exploitability eXchange) data to know if a CVE affects us, and then prioritizing direct dependency fixes ¹⁴ . For example, if a high-severity CVE is found in a direct library and it's reachable, it becomes a top-priority issue; if it's only in a rarely-used transitive dependency, it's lower priority ¹⁴ . All SCA findings (component, version, CVE, reachability) feed into our dashboard just like SAST findings.

In summary, the SCA module will (a) auto-generate SBOMs for each codebase ⁹ , (b) continuously update vulnerability information and license data for each SBOM component, and (c) perform intelligent triage by considering exploitability/reachability ¹² ¹⁴ . This follows best practices of enterprise SCA, ensuring we not only know *what* is in use but which of those issues truly threaten security.

Security, Compliance, and Audit

Robust security controls and auditability are built in from the start. We will implement **Role-Based Access Control (RBAC)** with fine-grained roles (e.g. system admin, team admin, developer, auditor) and multi-tenancy. Access rights (which projects or functions a user can see) are tied to organizational units, so large enterprises can enforce least privilege. We will integrate with corporate identity providers (LDAP/AD, SAML/OIDC) so that user management and SSO comply with enterprise policy. As best practice, we separate policy decision (PDP) from policy enforcement (PEP) – for example using Open Policy Agent as a standalone PDP that scales independently ¹⁵ . This allows us to update authorization policies (e.g. new roles or rules) without redeploying application code.

Audit logging is mandatory. Every significant action (scan run, findings viewed, policies changed, login) is logged with timestamp, user ID, and context. Logs will be forwarded to a central SIEM or ELK stack. Qualys recommends sending host and app logs (logins, privilege changes, data transfers, etc.) to a central collection platform ¹⁶ ; we will do likewise for our platform. We will also capture immutable scan results (so a finding record includes the raw SARIF output) to ensure findings can be traced in audits. Encryption is enforced: TLS 1.3 for all network traffic; database fields like tokens or secrets encrypted at rest.

For compliance, we align with standards (ISO 27001, SOC2, NIST). For example, we adopt a "policy-as-code" approach as described by Palo Alto: encoding security policies in CI/CD checks so that, say, any deployment is automatically blocked if the SBOM contains known critical CVEs ¹⁷ . We will provide out-of-the-box compliance reports (e.g. OWASP Top 10, SCA license compliance, regulatory frameworks) and support audit trails of all activities. Regular third-party penetration tests and code audits of our platform will be part of our roadmap.

Hybrid Deployment Patterns

The platform will be **hybrid-friendly** by design. In a typical pattern, the core services (dashboard, orchestration, database) can run as a SaaS or in a customer-managed cloud, while local scan agents execute on-premises. For **cloud mode**, all components run in a managed K8s cluster. In **on-prem mode**, customers

deploy the same containers in their environment, pointed at a central SaaS policy manager or fully isolated. To accommodate strict networks, we will support an **agent-based data plane**: lightweight agents installed on developer machines or build servers that handle local scan tasks and only send sanitized results (no proprietary code) back to the control plane.

In practice, we will provide deployment options: a Helm chart or installer for fully on-prem deployments, and a hosted cloud offering for SaaS customers. Checkmarx notes that cloud scanning “designed to scale on demand” can run scans globally in parallel, whereas on-prem may need manual hardware upgrades ¹⁸. We will mitigate on-prem scaling issues by containerizing scans (so additional pods can be added) and by allowing hybrid: organizations can burst to our cloud, or run local agents that connect to the cloud when needed.

Following Qualys’s advice for hybrid environments, we will use host-based agents for scanning and centralized logging. For example, continuous vulnerability scanning in ephemeral containers or VMs can be done via agents, with reports sent to the central server ¹⁹. All logs from on-prem agents (scan logs, errors) will feed back to the central logging system. In essence, wherever the code lives (cloud repo, local repo, air-gapped network), an appropriate agent or API integration will enable scanning, so that hybrid and on-prem environments get consistent coverage. We will ensure the on-prem deployment supports air-gapped operation: the entire platform can run without Internet if needed, with updates delivered via signed packages.

Integration Strategy

We will tightly integrate into developers’ and security teams’ workflows. The platform offers:

- **CI/CD integration**: Plugins or CLI tools for Jenkins, GitLab CI, GitHub Actions, Azure DevOps, etc., so that scans automatically run on pull requests, merges, or nightly builds. Inline gating is supported: e.g. a pipeline can fail if a “blocker” severity issue is found. As Kong’s security guide recommends, SAST/SCA should be executed early (“shift left”) in the pipeline ²⁰ and builds should be rejected on critical findings.
- **SCM integration**: Connectors for GitHub, GitLab, Bitbucket to trigger scans on commits/PRs and to push results back as status checks or comments. We also integrate with container registries: after a Docker image build, an agent can extract the SBOM (via Syft) and push it to our service for SCA scanning.
- **Issue trackers and notifications**: Discovered vulnerabilities can automatically create tickets or issues in Jira, GitHub Issues, Azure Boards, etc., with drill-down details and remediation advice. We will also support sending alerts to email, Slack/MS Teams, or webhook endpoints.
- **IDE integration (longer-term)**: An IDE plugin (e.g. for VS Code or IntelliJ) could allow developers to see scan results or AI suggestions in-context as they code, similar to modern developer-centric tools ²¹. For MVP, we focus on CLI/CI integration but design APIs to allow future IDE hooks.
- **Policy as Code**: We expose policy APIs (e.g. Open Policy Agent bundles) so organizations can enforce custom rules (for example, “no dependency with license X” or “no code with critical vulnerabilities in production branches”). These policies integrate into pipeline checks ¹⁷.

By embedding security tools into CI/CD and SCM, we ensure minimal friction. For example, a GitHub PR can trigger an SAST + SCA scan and post comments on the PR with findings, enabling immediate triage. This direct integration into developer workflows (rather than a separate security portal) follows the best practice of developer-first AppSec ²¹. All integration points are configurable via an administration console or Infrastructure-as-Code templates, letting enterprises plug the platform into their existing toolchain seamlessly.

MVP Roadmap (0-6 months)

Month 1-2 – Core Platform & SAST: Establish the basic microservices framework. Deploy a proof-of-concept with one language (e.g. Python) and one SAST scanner (Semgrep) in a Docker container. Build the orchestration service and queue system; set up a simple UI to trigger scans and view raw results. Implement SARIF normalization ⁷ and a basic user database with RBAC. Integrate with one CI system (e.g. GitHub Actions) for pull-request scanning.

Month 3-4 – Expand Language Support & SCA: Add additional languages and scanners (e.g. Java with SpotBugs/CodeQL, JavaScript with ESLint rules). Improve the UI/dashboard (e.g. grouping findings by project). Introduce SBOM generation (Syft + CycloneDX) and vulnerability lookup (e.g. using OSS Index or NVD) ⁹ ¹¹. Begin implementing SBOM storage and a simple dependency graph viewer. Integrate basic SCA alerts.

Month 5 – Security Features & Integrations: Harden security: finalize RBAC roles, audit logging, and encrypt communications. Add rich scan results (CVE details, remediation links). Integrate with Jira or GitHub Issues for auto-ticketing. Extend CI integrations (Jenkins, GitLab). Deploy on a public cloud environment for SaaS trials. Conduct internal compliance checks.

Month 6 – Refinement & Scaling: Optimize performance (parallel scans, caching). Implement findings deduplication and alert thresholds. Add more languages/frameworks (C/C++, C#, Go). Pilot the platform with a test customer or open beta. Gather feedback on the user experience and iterate. By end of month 6, have a minimally viable AppSec platform that provides automated SAST/SCA, a web UI, hybrid deployment options, and integrations, ready for broader roll-out.

Comparison to Similar Tools

Our platform combines and extends features of market tools like Xbow, Snyk, and Checkmarx:

- **Xbow** – An AI-driven penetration testing platform. Xbow uses autonomous AI agents to “discover, validate, and exploit vulnerabilities” continuously ²². It targets DevOps with “agentic AI” for full-stack pentesting. In contrast, our initial focus is on *static* (SAST/SCA) analysis. We plan to add dynamic and AI-driven pentesting (DAST/IAST) later, but will leverage lessons from Xbow’s reception. For example, users have noted Xbow’s limited developer workflow integration and US-only hosting as issues ²³. We will design for multi-region deployment (to meet GDPR/compliance needs) and deep IDE/CI integration from the start, avoiding those pitfalls. We also aim to make our output transparent and explainable – a common critique was that Xbow’s results were opaque and high in false positives ²⁴, so we will include context like reachability analysis ¹² and clear evidence for each finding.
- **Snyk** – A developer-focused SCA platform (also offering SAST plugins). Snyk excels at scanning open-source dependencies and integrates well into dev workflows. However, Checkmarx observes that “Snyk is just an SCA tool” and lacks the breadth of source-code analysis ²⁵. Our blueprint covers both SCA and full SAST, giving end-to-end coverage. We’ll match Snyk’s developer-friendliness (e.g. we target IDE plugins in the future) while also catering to enterprise needs (extensive language support, on-prem options). According to Checkmarx, solutions that scan only at commit time can

lead to delays and false positives, whereas inline IDE scanning finds issues earlier ²⁶. We aim to provide timely feedback (fast scans in CI) and eventually real-time guidance (e.g. pre-commit or IDE), blending Snyk's ease-of-use with a deeper analysis engine.

- **Checkmarx (One)** – A leading enterprise SAST solution. Checkmarx supports 80+ languages and prioritizes scale and customization. Like Checkmarx, we plan multi-language support and enterprise RBAC. We differ by starting in the cloud with hybrid flexibility and by deeply integrating AI. Checkmarx's newer "Developer Assist" emphasizes keystroke-level scanning and AI fixes ²¹ ²⁶. We will incorporate similar ideas: shift-left scanning and AI-suggested remediation, but as optional plugins. In short, our platform will marry Checkmarx's breadth with Snyk's agility and add an AI layer inspired by XBOW's vision.

Extensibility Roadmap

The platform is designed to grow beyond SAST/SCA:

- **DAST (Dynamic Testing)**: In a later phase, we will integrate or develop runtime scanners (e.g. OWASP ZAP, Burp) to test running applications. DAST can find issues (like SQLi or XSS) that static analysis misses. We'll orchestrate DAST jobs in staging environments and feed results into the same dashboard.
- **IAST (Interactive Testing)**: We plan to offer an agent that instruments applications during automated tests. According to Palo Alto Networks, IAST "combines insights from SAST and DAST in real time, running within the application to detect vulnerabilities as code executes" ²⁷. We could add libraries or sidecars that collect detailed code execution traces and flag issues on-the-fly, providing feedback even during QA or load testing.
- **Runtime Application Self-Protection (RASP)**: Eventually, we may include runtime hooks or a service mesh integration to block or report exploitation attempts in production. This extends protection beyond testing into live defenses.
- **Fuzzing and Bug Bounty Integration**: We may incorporate fuzz-testing tools or connect to bug-bounty programs via APIs to complement coverage.
- **Custom Security Rules and Machine Learning**: Over time, allow users to upload custom SAST/SCA rules or train ML models on their own codebase patterns.
- **DevSecOps Automation**: We will enhance the policy engine (OPA) with templates (e.g. PCI, HIPAA) and automate remediation workflows (ticketing, patch orchestration).
- **Application Security Posture Management (ASPM)**: Build a unified risk dashboard that correlates SAST/SCA/DAST/IAST findings across all apps, helping CISOs prioritize which teams or applications need urgent attention.

Each new capability will plug into the existing microservices framework and reuse core components (jobs queue, database, UI), ensuring the platform's value compounds. By designing modularly from the start, we can evolve continuously to cover new AppSec domains (like secrets detection, infrastructure as code scanning, container security, etc.) as enterprise needs grow.

Sources: Architecture and microservices guidance ¹; SAST/SCA best practices and SBOM guidelines ⁷ ⁹ ¹¹; AI in security (GPT-4 vs SAST) ⁵ ⁶; fine-tuning LLMs for code ⁴; policy and pipeline security ¹⁷ ¹⁹; hybrid/cloud-vs-onprem tradeoffs ¹⁸ ²⁸; tool comparisons (Checkmarx vs Snyk vs XBOW) ²⁵ ²⁶ ²³ ²⁴.

1 Microservices Architecture | Atlassian

<https://www.atlassian.com/microservices/microservices-architecture>

2 3 Best languages for microservices | Chronosphere

<https://chronosphere.io/learn/best-languages-for-microservices/>

4 Beating GPT-4: Fine-Tuning Llama-3 for Software Security - inovex GmbH

<https://www.inovex.de/de/blog/beating-gpt-4-fine-tuning-llama-3-for-software-security/>

5 6 GPT-4 Significantly Outperforms Static Analysis In Software Vulnerability Detection.

<https://quantumzeitgeist.com/gpt-4-significantly-outperforms-static-analysis-in-software-vulnerability-detection/>

7 The complete guide to SARIF: Standardizing static analysis results | Sonar

<https://www.sonarsource.com/resources/library/sarif/>

8 18 28 On-Prem vs. Cloud SAST: What Security Leaders Need to Know Before They Migrate

<https://checkmarx.com/learn/sast/on-prem-vs-cloud-sast-what-security-leaders-need-to-know-before-they-migrate/>

9 10 11 14 Dependency Graph SBOM - OWASP Cheat Sheet Series

https://cheatsheetseries.owasp.org/cheatsheets/Dependency_Graph_SBOM_Cheat_Sheet.html

12 13 Reachability Analysis in SCA - Labrador Labs

<https://labradorlabs.ai/news/reachability-analysis-in-sca/>

15 Best Practices for Microservice Authorization

<https://www.permit.io/blog/best-practices-for-authorization-in-microservices>

16 19 Securing the Hybrid Cloud: A Guide to Using Security Controls, Tools and Automation | Qualys

<https://blog.qualys.com/product-tech/2018/05/15/securing-the-hybrid-cloud-a-guide-to-using-security-controls-tools-and-automation>

17 27 What Is AppSec? - Palo Alto Networks

<https://www.paloaltonetworks.com/cyberpedia/appsec-application-security>

20 10 Microservices Security Challenges & Solutions for 2025 | Kong Inc.

<https://konghq.com/blog/engineering/10-ways-microservices-create-new-security-challenges>

21 25 26 Checkmarx vs Snyk: Why Choose Snyk Alternatives and Competitors?

<https://checkmarx.com/snyk/>

22 23 24 Top XBOW Alternatives In 2026

<https://www.aikido.dev/blog/xbow-alternatives>